

Intro to Structured Streaming

Why do we need stream processing?

- Streaming data can come in faster than it can be consumed when using traditional batch-related processing techniques.
 - Bank transactions
 - IoT devices
 - Online games
 - ...
- A stream of data is treated as a table to which data is continuously appended.

High level components of a streaming system

- Input source
 - Kafka
 - **Azure Event Hubs**
 - IoT Hub
 - Network sockets
 - ...
- Stream processing
 - **Structured Streaming**
 - foreach sinks
 - Memory sinks
 - ...

Azure Databricks structured streaming

- Fast, scalable, fault tolerant stream processing API
- **Near** real-time analytics on streaming data
- Processing streaming data is very similar to processing static data
- The API continuously increments and updates the final data
- Can work with Azure Event Hubs and analyze data using structured streaming query and Spark SQL

Spark uses a single API to handle batch and streaming data..

Batch ETL with DataFrames

```
input = spark.read  
    .format("json")  
    .load("source-path")
```

Read from Json file

```
result = input  
    .select("device", "signal")  
    .where("signal > 15")
```

Select some devices

```
result.write  
    .format("parquet")  
    .save("dest-path")
```

Write to parquet file

Streaming ETL with DataFrames

```
input = spark.read  
    .format("json")  
    .stream("source-path")
```

Read from Json **file stream**
Replace **load()** with **stream()**

```
result = input  
    .select("device", "signal")  
    .where("signal > 15")
```

Select some devices
Code does not change

```
result.write  
    .format("parquet")  
    .startStream("dest-path")
```

Write to Parquet **file stream**
Replace **save()** with **startStream()**

- `read...stream()` creates a streaming dataframe, but does not start any of the computation
- `write...startStream()` defines where and how to output the data and starts the processing
 - Basically a query that runs repeatedly, updating the output each time


Append mode

- This is a simple query without any aggregation..

```
input = spark.read
    .format("json")
    .stream("source-path")

result = input
    .select("device", "signal")
    .where("signal > 15")

result.write
    .format("parquet")
    .startStream("dest-path")
```



What if we have Aggregations?

- Continuous windowed aggregations

```
input.avg("signal")
```

Continuously compute *average*
signal *across all devices*

```
input.groupBy("device-type")  
  .avg("signal")
```

Continuously compute *average*
signal of *each type of device*

- Continuous windowed aggregations

```
input.groupBy(  
  $"device-type",  
  window($"event-time-col", "10 min"))  
  .avg("signal")
```

Continuously compute
average signal of *each type*
of device in last 10 minutes
using *event-time*

Output Modes

- Different output modes make sense for different queries
- **Append mode for non-aggregation queries**

```
input.select("device", "signal")  
  .write  
  .outputMode("append")  
  .format("parquet")  
  .startStream("dest-path")
```

- **Complete mode for aggregation queries**

```
input.agg(count("*"))  
  .write  
  .outputMode("complete")  
  .format("parquet")  
  .startStream("dest-path")
```

Query Management

- Multiple queries can be active at the same time
- We can assign a query to a variable and manage it. We can:
 - stop execution, wait for it to terminate
 - Get its status
 - Get the exception if there is an error

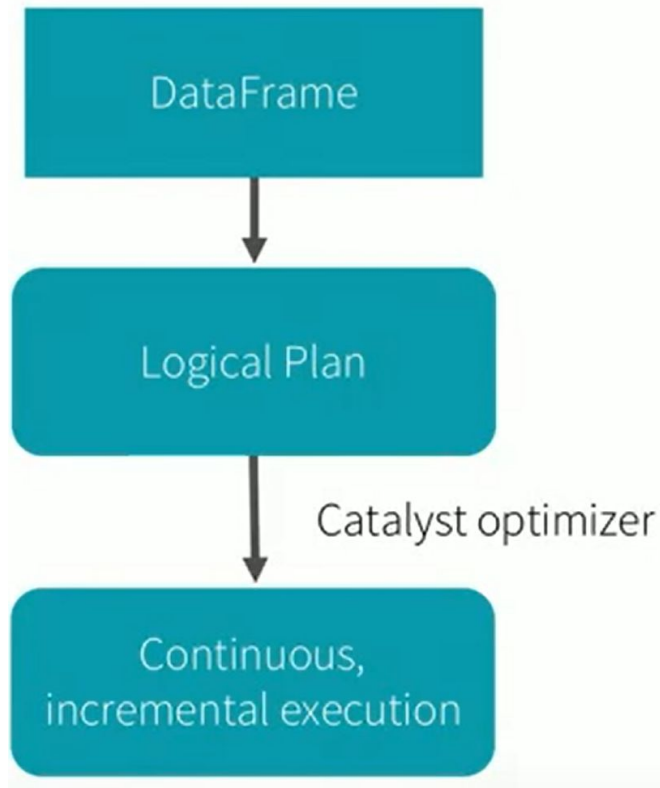
```
query = result.write  
    .format("parquet")  
    .outputMode("append")  
    .startStream("dest-path")
```

```
query.stop()  
query.awaitTermination()  
query.exception()
```

```
query.sourceStatuses()  
query.sinkStatus()
```

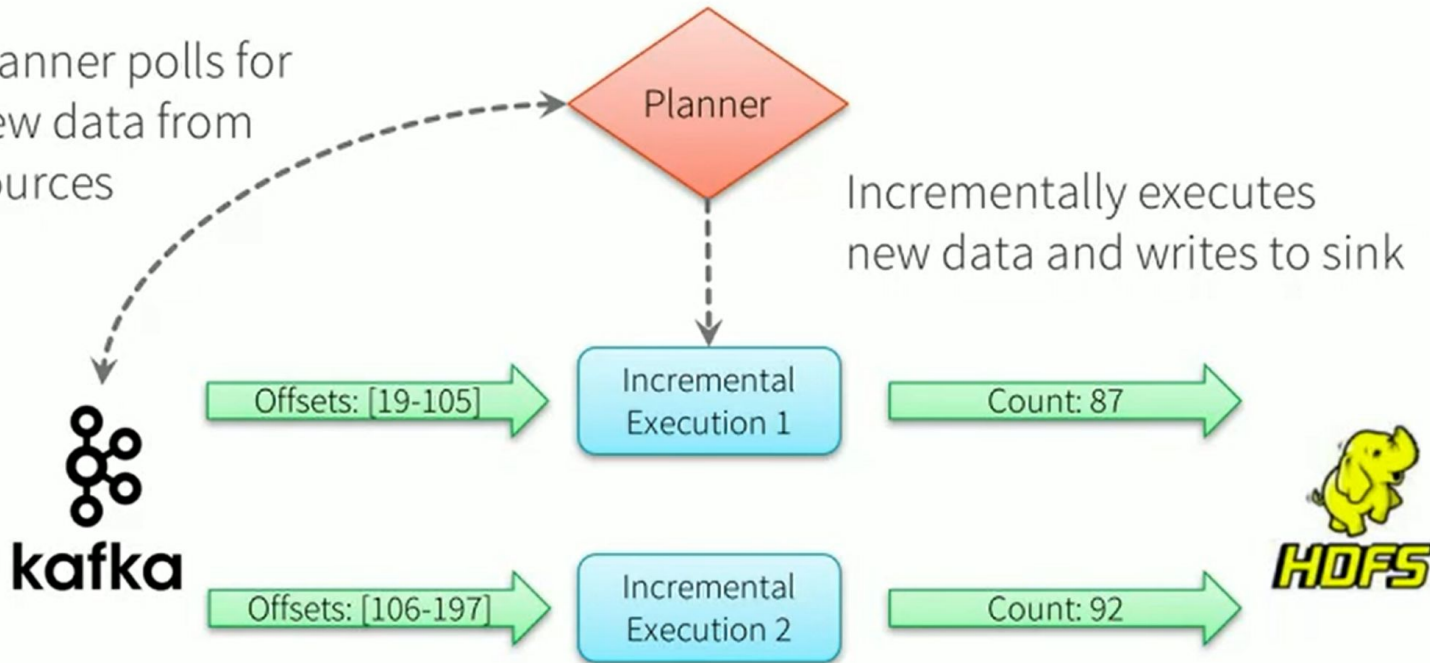
How are the queries executed?

- **Logically**, we are writing queries as if we are working on a static table
- **Physically**, Spark automatically runs the query *incrementally* and *continuously*



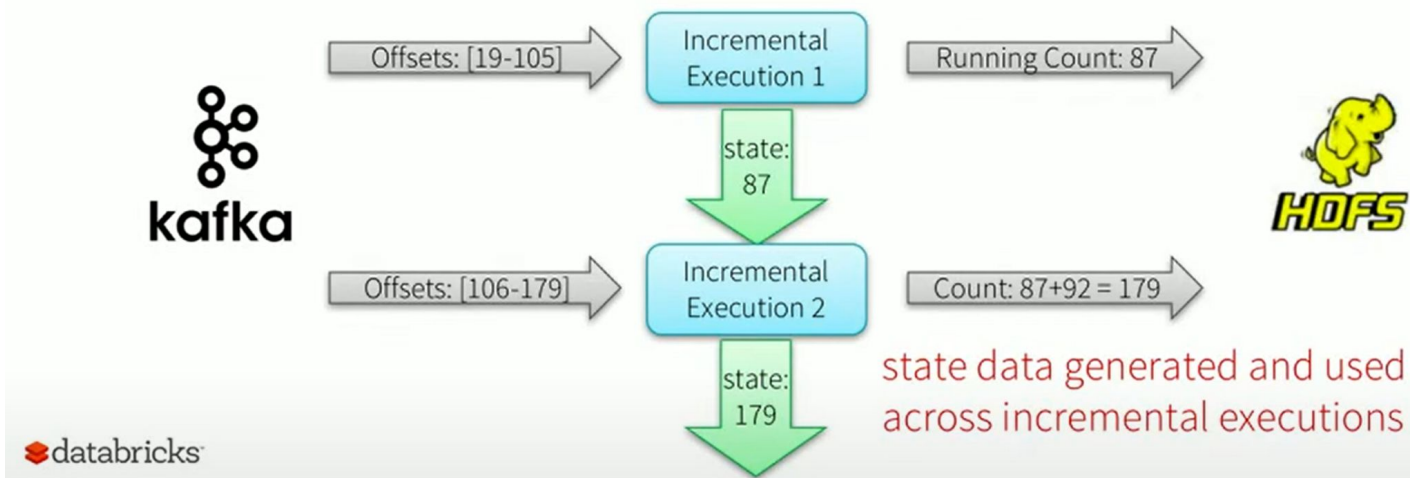
How are the queries executed?

Planner polls for
new data from
sources



How are the queries executed?

- Previous result is stored in memory and used for the next execution



How is fault-tolerance achieved?

- Everytime the planner gets a offset of data, the offset is first stored before any execution
 - If execution fails, we can read back the offset and rerun the execution
 - This relies on the datasource to be able to replay the data based on the offset
 - By design, structured streaming sources are replayable(Kafka, kinesis, file...)
- Planner also makes sure to use the correct version of in-memory state for the execution(for aggregations)
- Sinks, by design, will make sure to handle re-executions correctly and avoid double committing the output.
- **Offset tracking + in-memory state management + fault-tolerant sources and sinks = end-to-end fault tolerance**

Azure Event Hubs

- A scalable real-time data ingestion service
- Can receive large amounts of data from multiple sources and stream the prepared data to ADLS or Blob storage
- Can be integrated with Spark Structured Streaming to perform stream processing

Creating Event Hubs

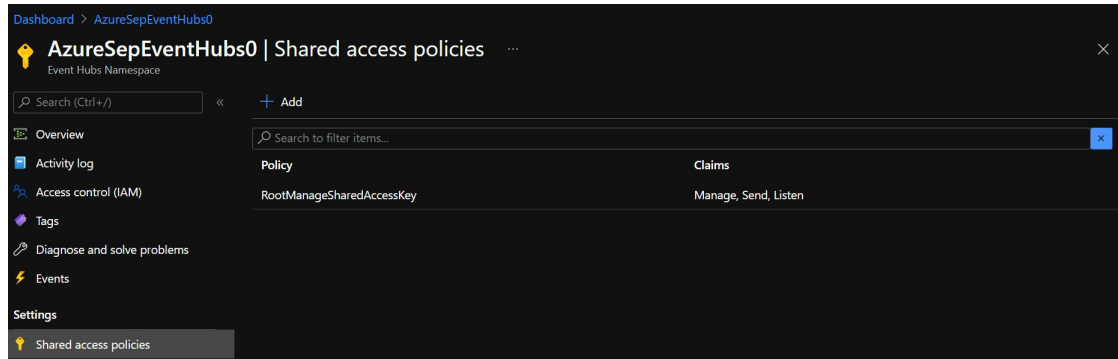
- Click “Create a resource” in the portal
- Search “Event Hubs” and create the resource
- Locate the service within the resource group you selected
- Click “+ Event Hub”

Settings in Event Hub creation

- Partition count
 - At least one partition per Throughput Unit(TU)
 - TUs can be scaled independently from the partition count
 - Cost is calculated based on the TUs only and not on the number of partitions
 - More partitions make the processing more complex, and it's only beneficial when you have tens of substreams.
- Message Retention
 - Locked to 1 day for Basic tier namespace
 - Up to 7 days for Standard tier
- Capture
 - Not available for Basic tier
 - Specify a minimum size and time window to perform the capture
 - Captured data can be automatically saved to a Blob or ADLS account of your choice
 - Data will be in Apache Avro format

Getting the connection string

- After the creation is finished, you can access the connection string from “Shared access policies”



Reading Streams

- **SparkSession.readStream** returns a **DataStreamReader** used to configure the stream
 - Recall **SparkSession.read**, which returns a **DataFrameReader** that can be used to read data in as a DataFrame.
- Key things to configure:
 - Schema
 - Predefined for you in some Pub/Sub sources like Kafka and Event Hubs
 - User-defined only(no inferencing) for file-based streaming
 - Stream type
 - Kafka
 - Files
 - TCP/IP
 - Stream type specific configurations

Reading Streams

- Creating a streaming df is very similar to creating a static df, except we have to specify the schema

Cmd 12

```
1 streaming_df = (spark
2                 .readStream
3                 .option("maxFilesPerTrigger",1)
4                 .schema(schema)
5                 .json(file_path)
6                 )
```

- We can check if a df is streaming with **df.isStreaming**, which returns **True** if the df is a streaming df(vs. A static df)

Operations on Streaming Dataframes

- Most operations are identical to operations on a static df
- Some operations are not supported
 - Sorting -> only supported after aggregation and in Complete Output Mode
 - Limit / take(n)
 - Distinct
 - Certain types of joins
 - ...
 - Check doc for more details:
<https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html#unsupported-operations>

Join types supported

Left Input	Right Input	Join Type	
Static	Static	All types	Supported, since its not on streaming data even though it can be present in a streaming query
Stream	Static	Inner	Supported, not stateful
		Left Outer	Supported, not stateful
		Right Outer	Not supported
		Full Outer	Not supported
		Left Semi	Supported, not stateful
Static	Stream	Inner	Supported, not stateful
		Left Outer	Not supported
		Right Outer	Supported, not stateful
		Full Outer	Not supported
		Left Semi	Not supported
Stream	Stream	Inner	Supported, optionally specify watermark on both sides + time constraints for state cleanup
		Left Outer	Conditionally supported, must specify watermark on right + time constraints for correct results, optionally specify watermark on left for all state cleanup
		Right Outer	Conditionally supported, must specify watermark on left + time constraints for correct results, optionally specify watermark on right for all state cleanup
		Full Outer	Conditionally supported, must specify watermark on one side + time constraints for correct results, optionally specify watermark on the other side for all state cleanup
		Left Semi	Conditionally supported, must specify watermark on right + time constraints for correct results, optionally specify watermark on left for all state cleanup

Writing a Stream

- **DataFrame.writeStream** returns a **DataStreamWriter** used to configure the output of the stream.