The big picture

Lesson goals

- Recall foundational concepts related to the contemporary big data landscape.
- Define the core challenges with building a big data architecture.
- Explain how the core components of a Lakehouse relate to the Inmon Architecture.
- Explain how Delta Lake can be used to build a Lakehouse.
- Describe the core components of Delta Lake.

The big data landscape

databricks

The big data problem

Volume

Velocity

Variety

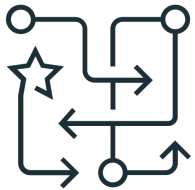

Veracity

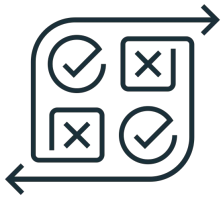

Value

At Moovio, your SLA includes:

Data freshness

Query speed

Data reliability

Ease of use

Big data needs

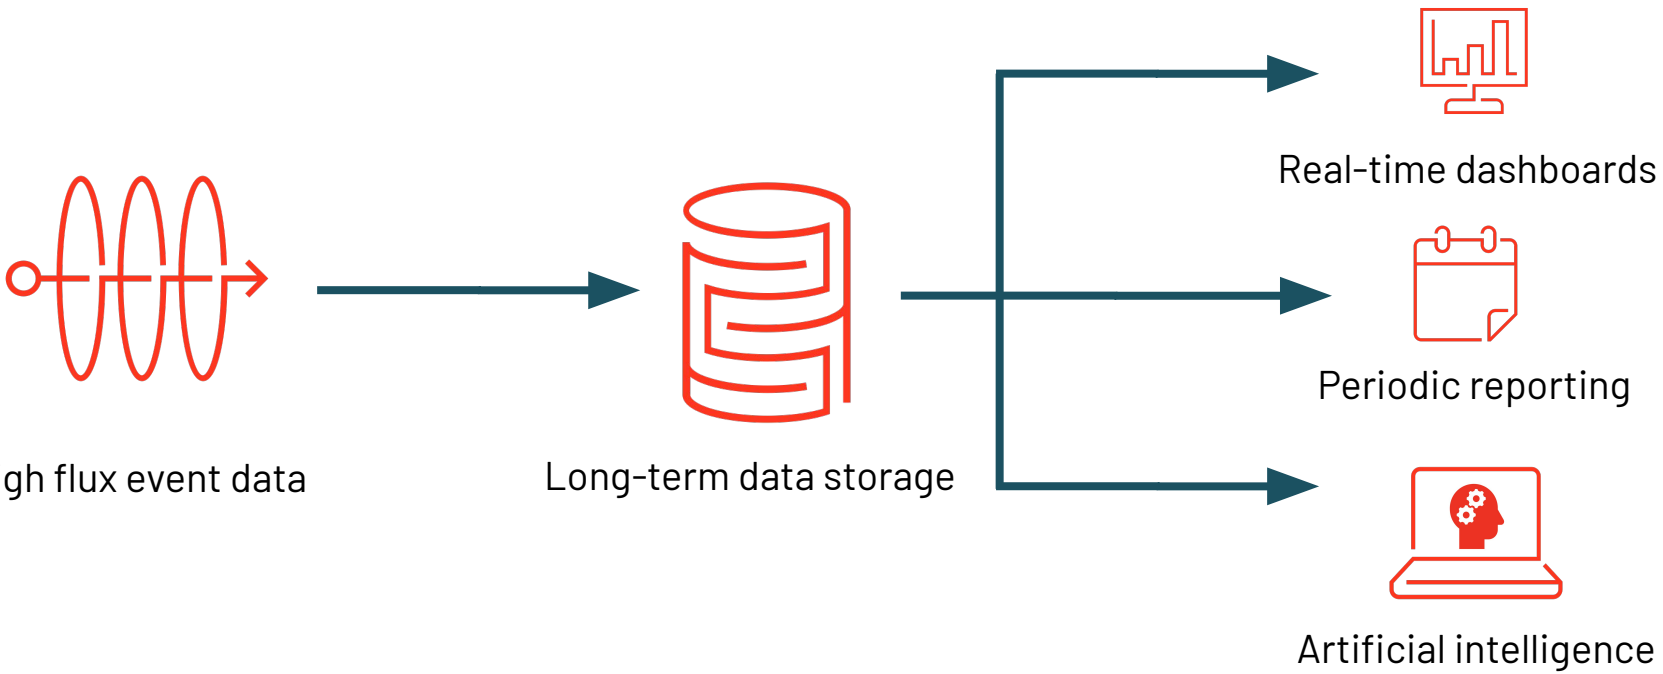

A single source of truth

ODS and OLTP

Operational Data Store and Online Transaction Processing

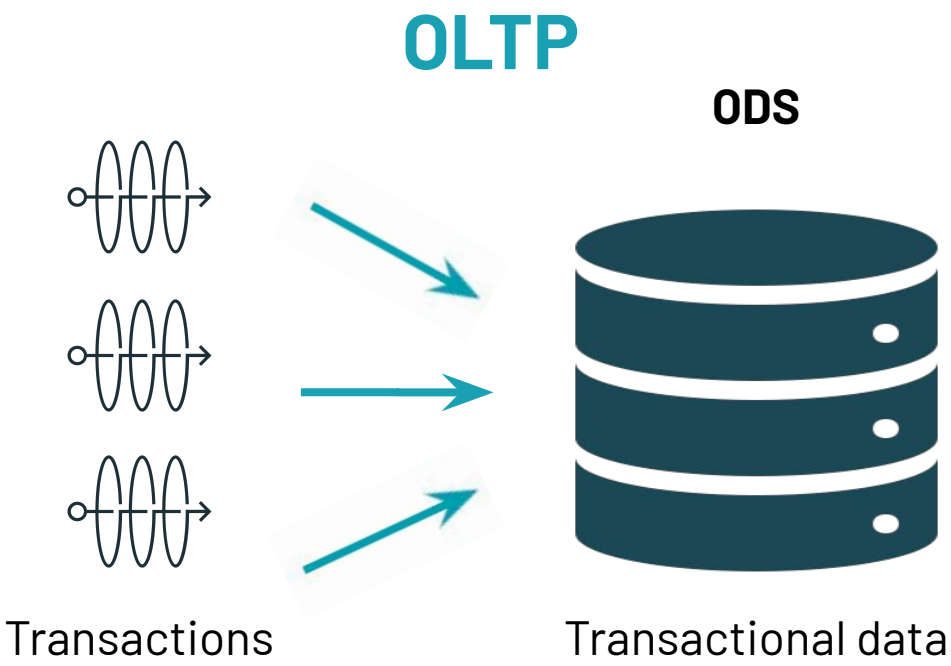

EDSS and OLAP

The Enterprise Decision Support System and Online Analytical Processing

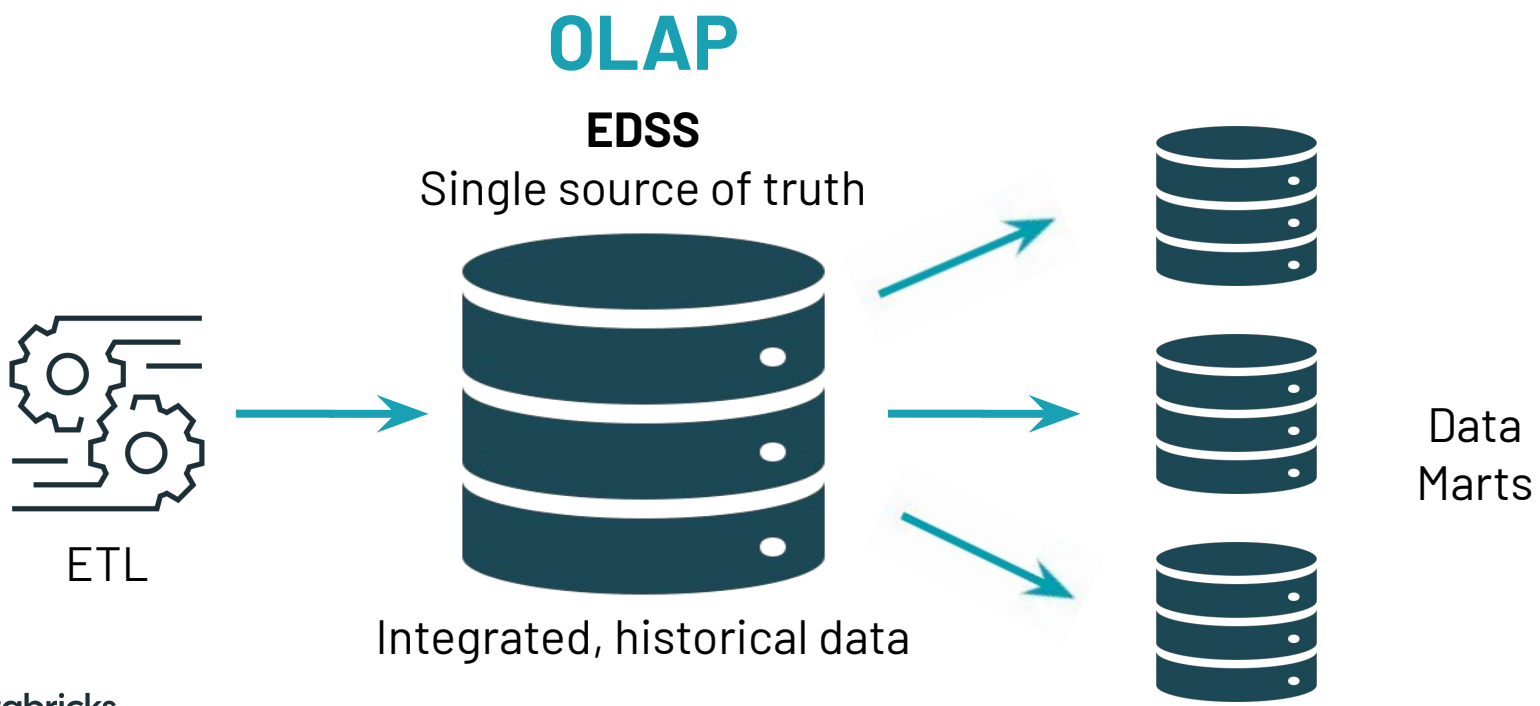

Complete data system

- An ETL process pulls data from the ODS to be loaded into the EDSS

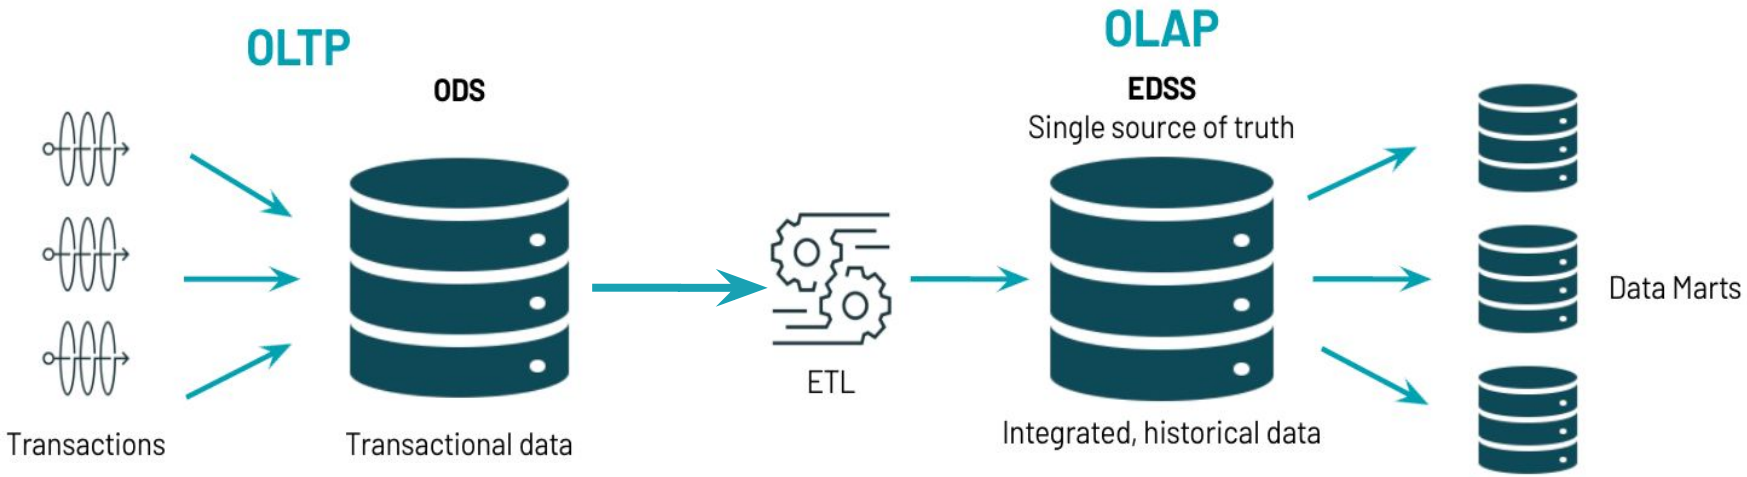

Levels of data

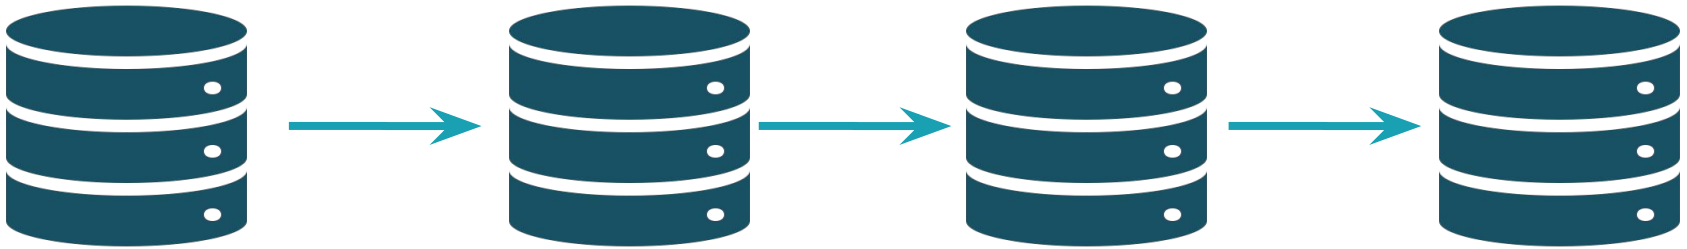

Fact and dimension tables

Dimensional modeling

- Fact tables
- Dimension tables
- Aggregate fact tables

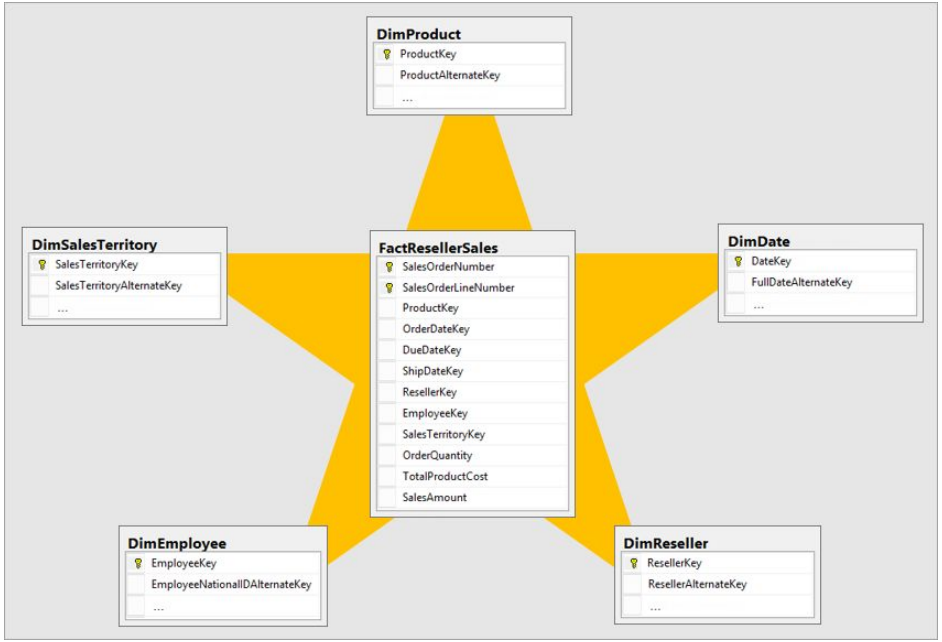

The Lakehouse

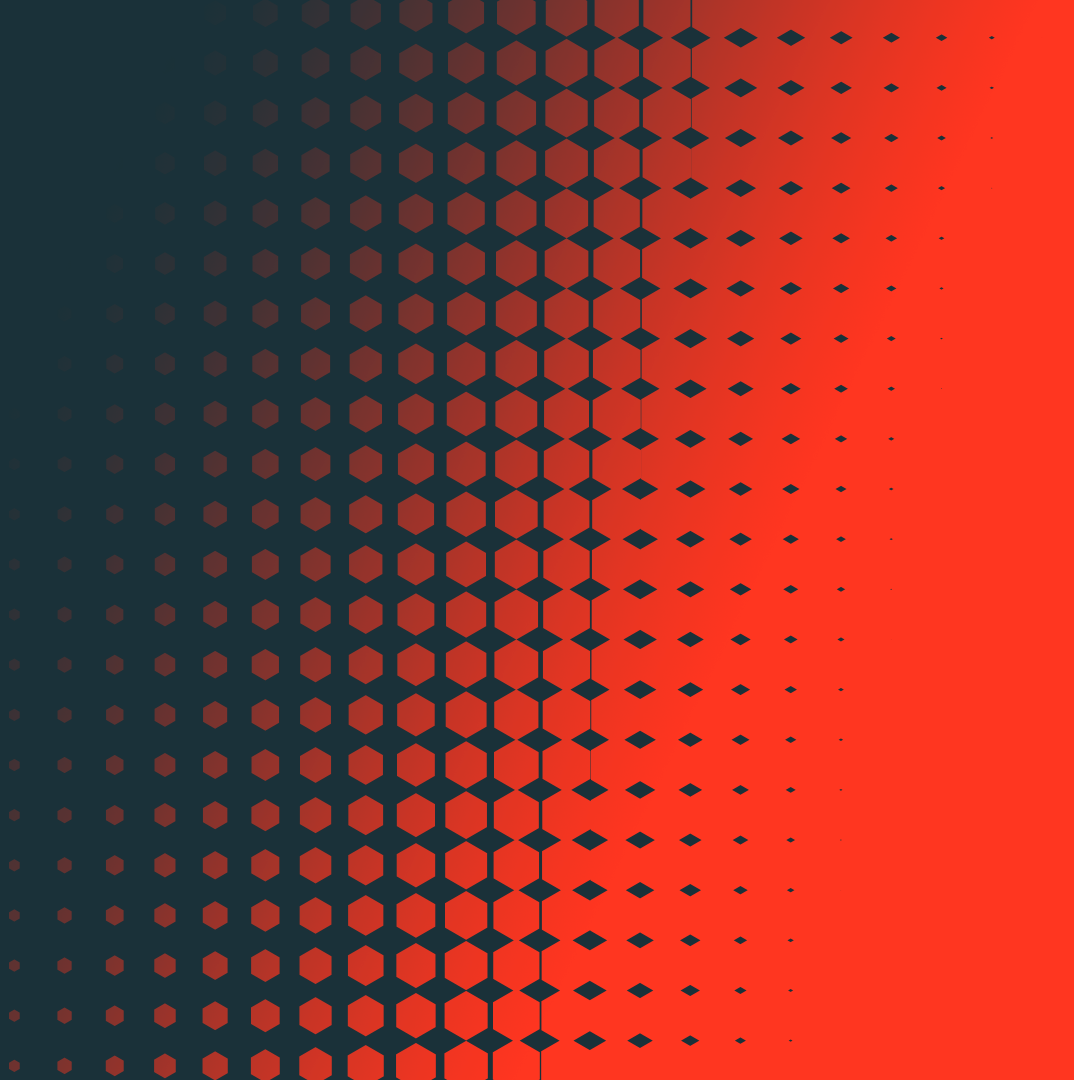

databricks

What is a Lakehouse?

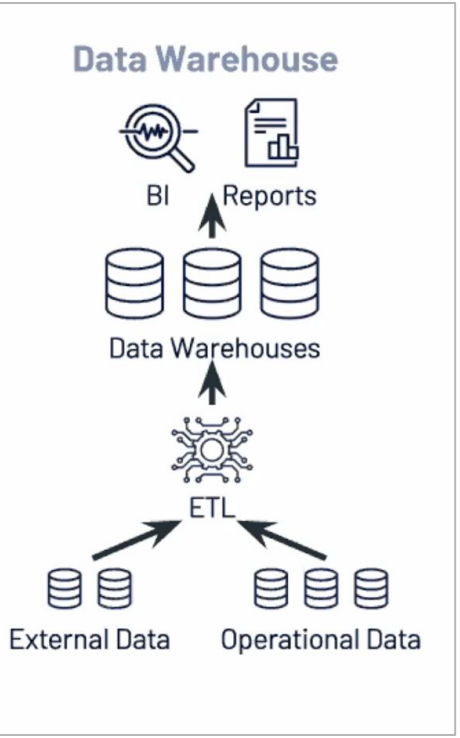

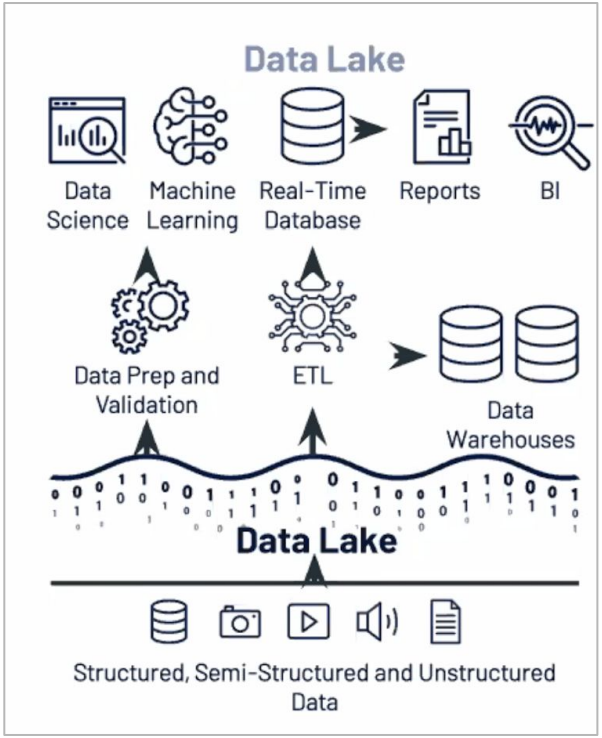

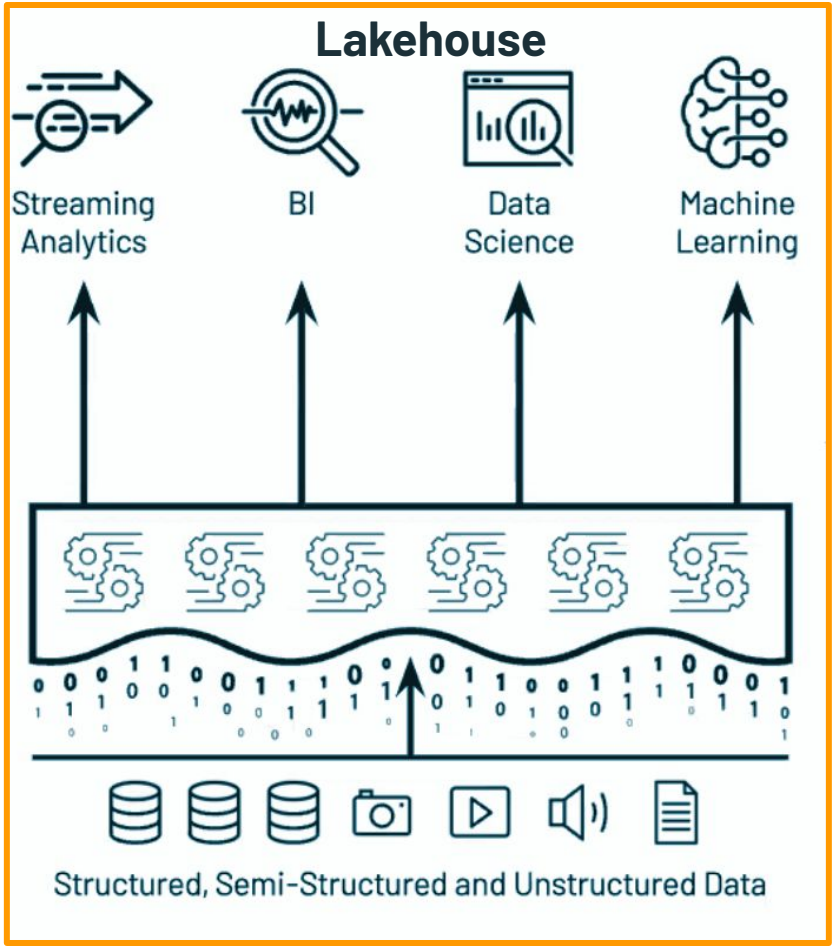

databricks

Levels of data in a Lakehouse

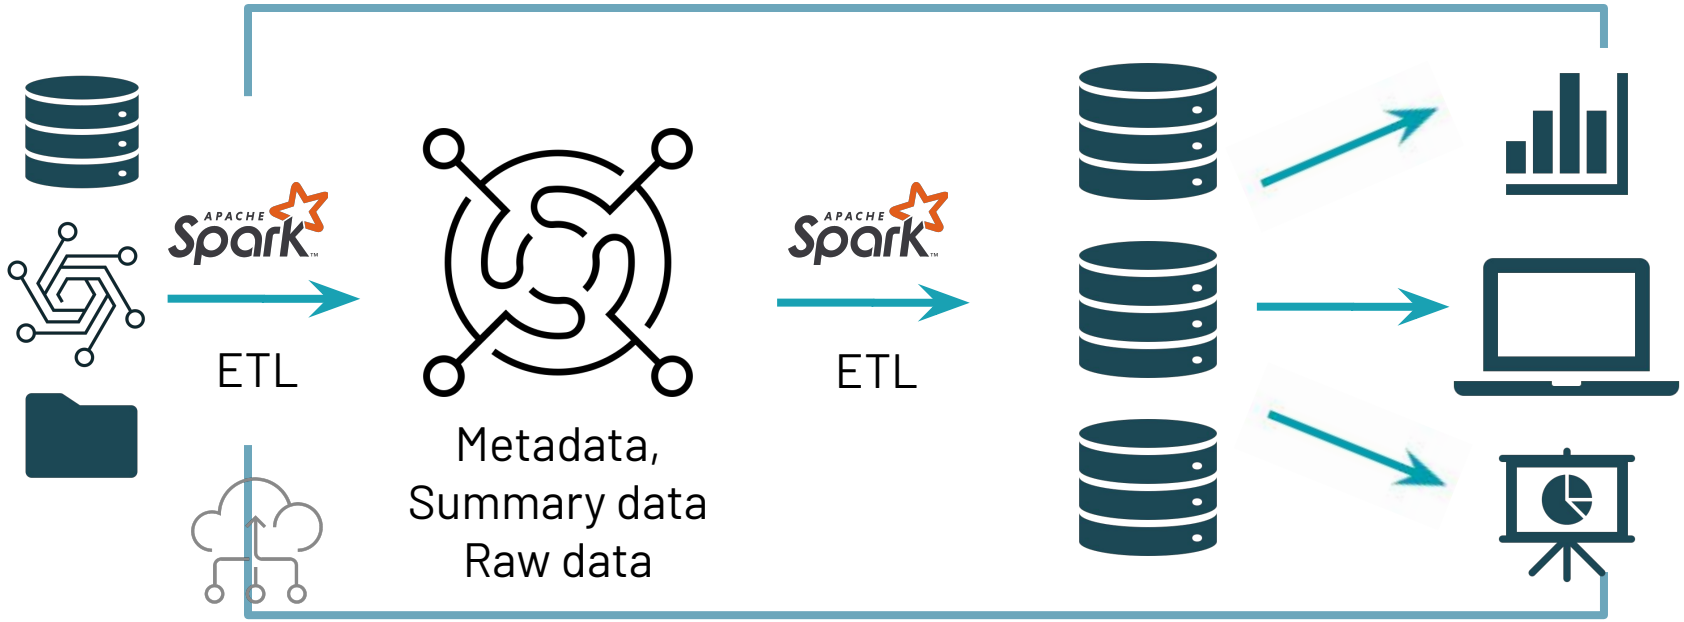

Operational
ODS

Atomic
EDSS
Data Lake

Departmental
Data Marts
Data warehouse

Individual
End users

databricks

Benefits of a Lakehouse

- Separation of compute and storage
- Infinite storage capacity
- Leverage best aspects of a data warehouse
- Low data gravity
- High data throughput
- No limits on data structure
- Mix batch and streaming workloads

The Delta architecture

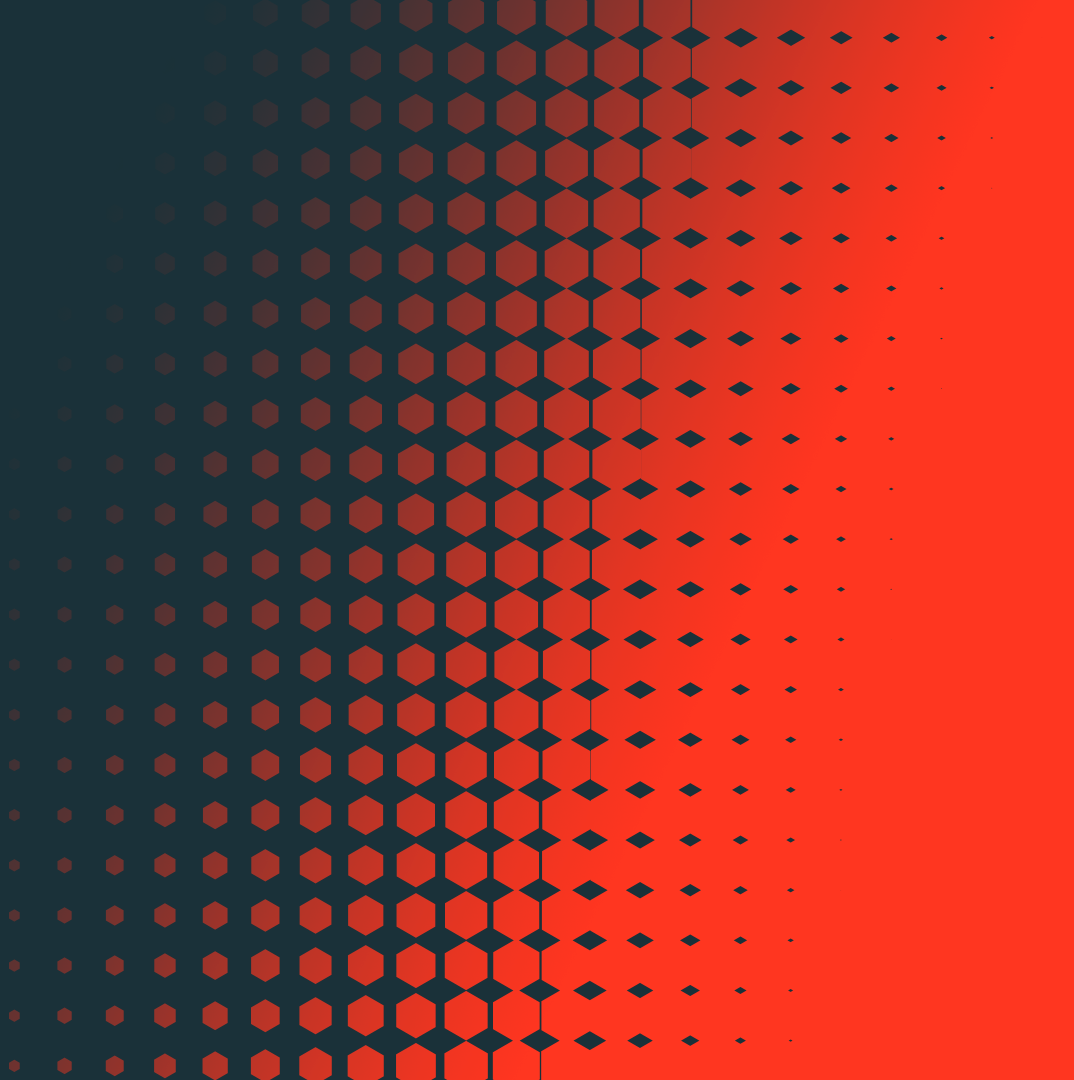

databricks

What is **Delta Lake**?

- Technology designed to be used with Apache Spark to build robust data lakes

Delta Lake features

- ACID transactions on Spark
- Scalable metadata handling
- Streaming and batch unification
- Schema enforcement
- Time travel
- Upserts and deletes
- Fully configurable/optimizable
- Structured streaming support

Delta Lake components

Delta Lake components

Delta Lake storage layer

- Highly performant and persistent
- Low-cost, easily scalable object storage
- Ensures consistency
- Allows for flexibility

Delta Lake components

Delta table components

- Data in Parquet/Delta files
- Transaction log
- Registered in metastore (optional)

Data - Parquet files

- File format for tabular data stored as columns
- Fast and powerful
- Delta files = Parquet + versioning + metadata

Transaction log

- Record of all transactions on a Delta table
- Prevents read conflicts
- Commits - ordered, atomic, json files
- Created automatically in the *_delta_log* subdirectory

Delta Lake components

Delta Engine

- File management optimizations
- Auto-optimized writes
- Performance optimization via Delta caching

The goal of a data engineer

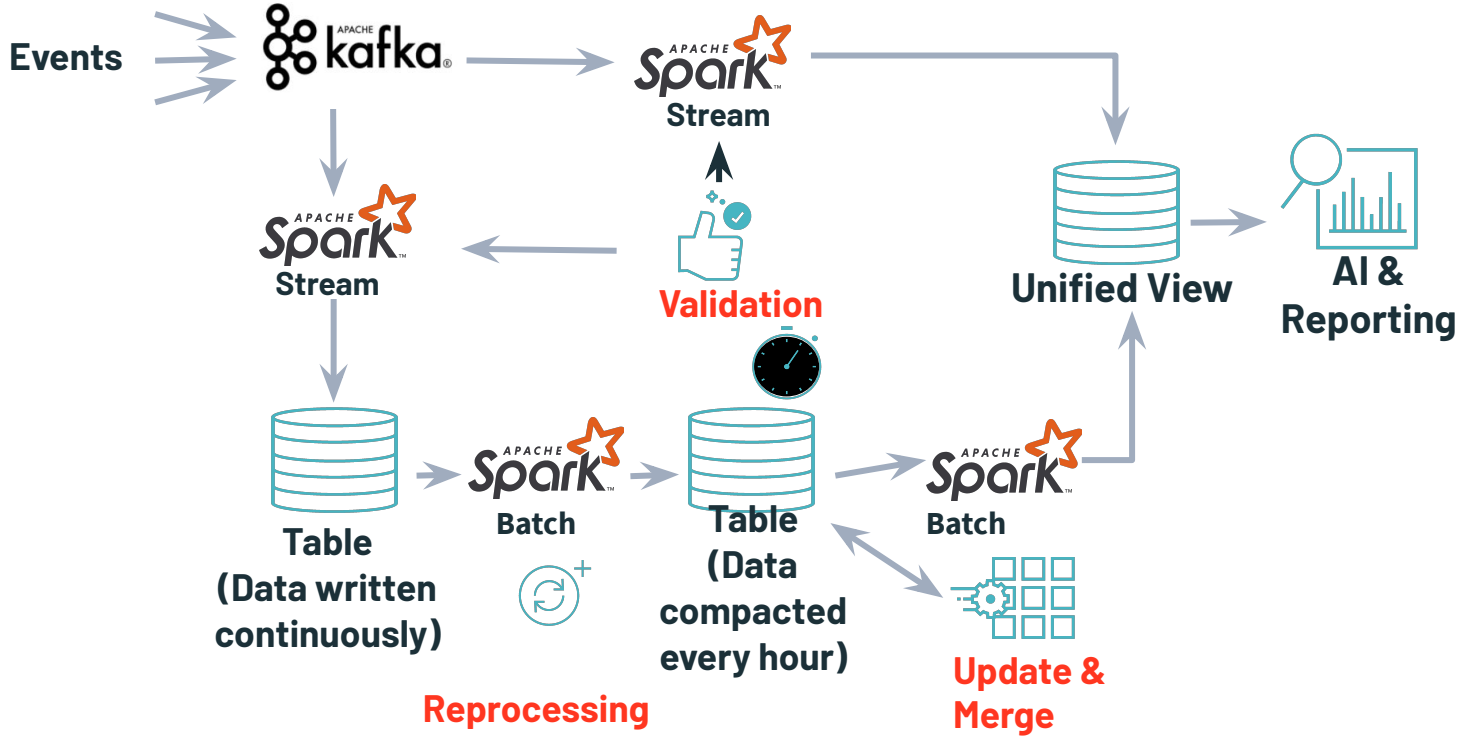

The Delta architecture design pattern

Delta architecture - Bronze

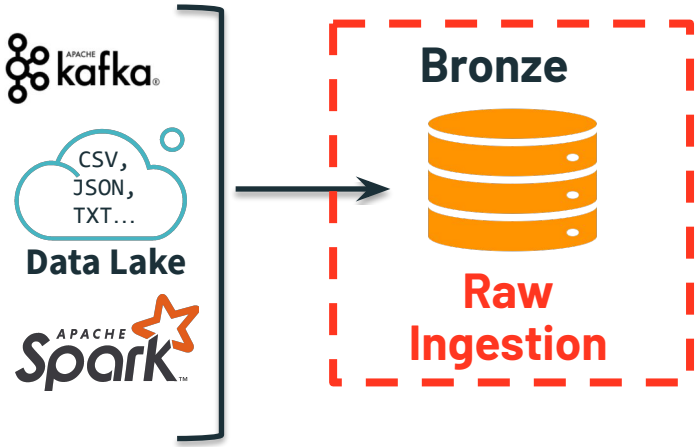

Delta architecture - Silver

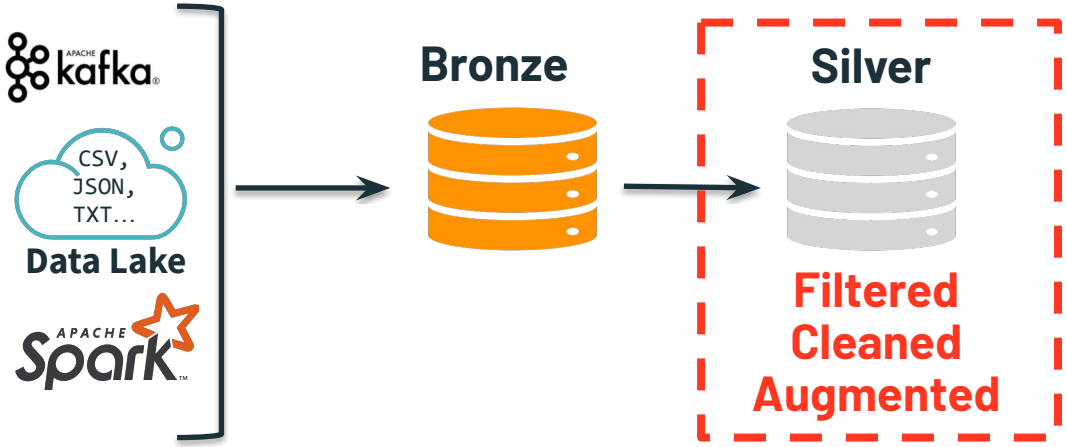

Delta architecture - Gold

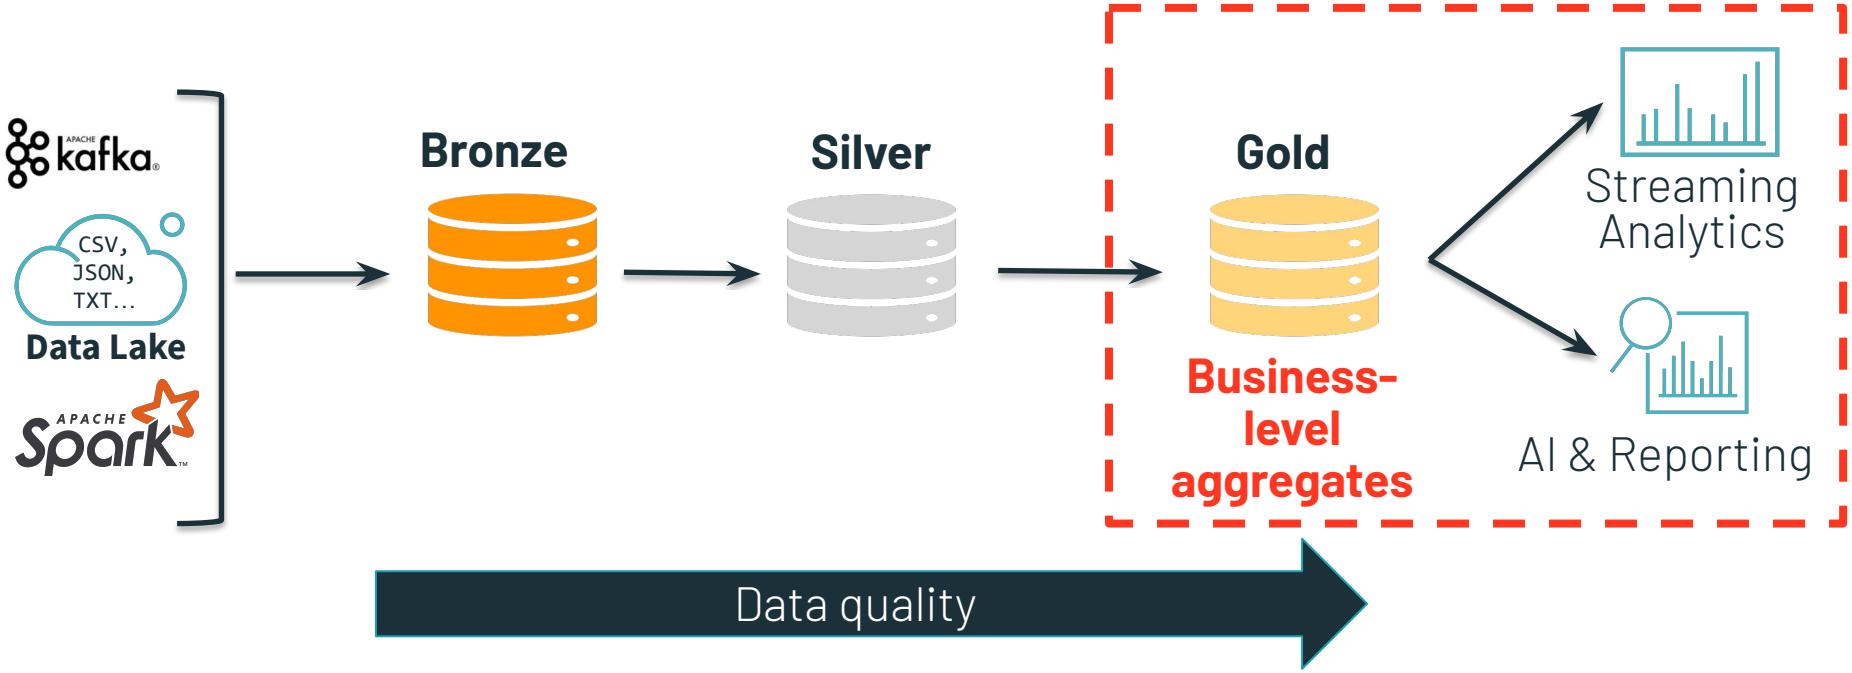


databricks